

input\_dataset — List of vectors or lists with integers representing words in review.  
Every individual vector /list represents a review.

target\_dataset — list of 0's and 1's  
negative. Positive  
review. review

images input pixels in  
 $\downarrow$   $\downarrow$  one image.  
Matrix  $(1000, 784)$  — gives one vector/ image

Weights\_0\_1  $(784, 40)$  — 40 (hidden layer and # of  
 $\uparrow$   $\uparrow$  neurons = 40  
rows columns

Weights\_1\_2  $(40, 10)$  — 10 (outputs)  
 $\uparrow$   $\uparrow$   
rows columns

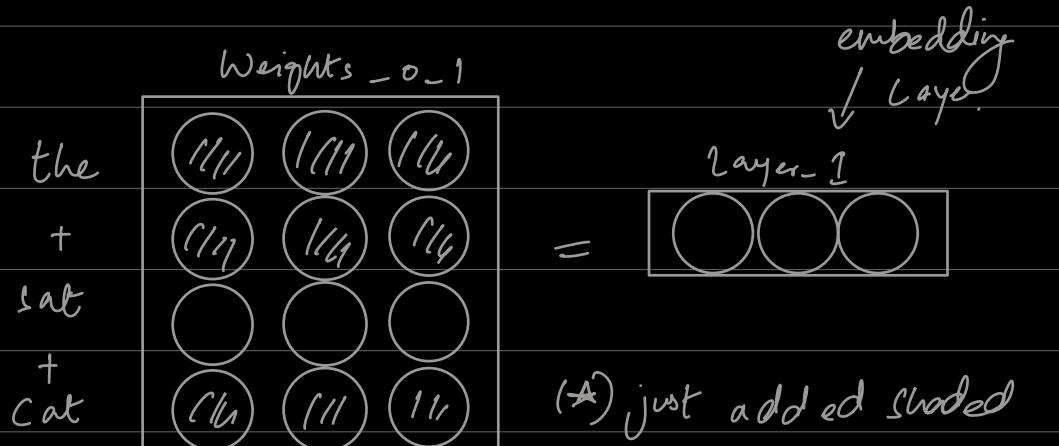
processing all images/vectors — 350 times to  
reduce error.

$$\left[ \begin{array}{c} \\ \\ \\ \\ \end{array} \right]$$
 } matrix of reviews and  
every vector /row  
represents a review and  
every review consisting

\*) weights<sub>-0-1</sub> and weight<sub>-1-2</sub>  
(Linear Layers) bunch of 0's and 1's  
made using one hot

\*) We can take a bit of shortcut  
to layer<sub>-1</sub> by replacing the  
first linear layer weights<sub>-0-1</sub> with  
an embedding layer. vocab or words in a  
review

\*\*\* ) Taking a vector of 1's and 0's is equivalent  
to summing several rows of a matrix. Instead  
of vector multiplication — sum relevant rows  
weights<sub>-0-1</sub>



(\*) just added shaded

rows rather than matrix

vector multiplication

$24,000$  - reviews  $\Rightarrow 24,000$  dimensions

Total vocab / words -  $70,000$

weights\_0\_1  $\rightarrow (70,000, 100)$

Weights\_0\_1  $\rightarrow (100, 1)$

$\underbrace{100 \text{ cols}}$



input\_dataset  $\rightarrow (24,000, 70,000)$   $\stackrel{\text{reviews}}{\downarrow} \stackrel{\text{cols}}{=}$

$24,000, 70,000$   
processing  $(1, 70,000)$  - vector / review

$(70,000, 100)$  - weight\_matrix

$$(1, 70,000) \circ (70,000, 100) = (1, 100)$$

vector  
of  
inputs / neurons

Matrix

vector

inputs  
!

weights  
!

$70,000 \cdot 70,000$   $\underbrace{?}_{\text{takes place 100 times since}}$

100 cols → and every

multiplication yields a  
neuron.

\* However, in this case we would not multiply  
and simply add rows of weight matrix corresponding  
one input vector of  $1 \times 70,000 \rightarrow$  giving us  $1 \times 100$   
vector or embedding layer.

\*  $[1 \times 100]$  vector then gets multiplied to  $100 \times 1$   
matrix or vector. — which yields an output of 1 or  
0 — where 1 represents positive.  
and 0 represents negative

⇒ Training neural network to predict if review is positive  
or negative.

```
import sys
```

```
f = open('reviews.txt')
```

```
raw_reviews = f.readlines()
```

```
f.close()
```

```
f = open('labels.txt')
raw_labels = f.readlines()
f.close()
```

```
tokens = list(map(lambda x: set(x.split(" ")),
                  raw_reviews)) # list of sets, sets
representing individual reviews.
```

*pre-processing*

```
Vocab = set
for sent in tokens: # iterating through individual set/review.
    for word in sent:
        if len(word) > 0:
            vocab.add(word)
Vocab = list(vocab) # all words
```

```
word2index = {} # word to integer/index
for i, word in enumerate(vocab):
    word2index[word] = i
```

```
input_dataset = list()
for sent in tokens:
```

```
    sent_indices = list()
    for word in sent:
```

try :

sent\_indices.append(word2index[word])

except:

" "

input\_dataset.append(list(set(sent\_indices)))

target\_dataset = List()

for label in raw\_labels:

if label == 'positive\n':

target\_dataset.append(1)

else:

target\_dataset.append(0)

# input\_dataset is List of vectors/lists with integers/

indices representing words in a review.

↳ a vector or list of words as integers/indices in  
a review.

```
import numpy as np  
np.random.seed(1)
```

```
def sigmoid(x):  
    return 1 / (1 + np.exp(-x))
```

alpha, iterations = (0.01, 2)

hidden\_size = 100 # features

weights\_0\_1 = 0.2 \* np.random.random((len(vocab), hidden\_size)) - 0.1

weights\_1\_2 = 0.2 \* np.random.random((hidden\_size, 2)) - 0.1

correct, total = (0, 0)

for epoch in range(iterations):

Trains on first 24,000 reviews.

for i in range(len(input\_dataset) - 1000):

# Loading a x, y = (input\_dataset[i], target\_dataset[i])

review and a layer\_1 = sigmoid(np.sum(weights\_0\_1[x], axis=0))

label.

layer\_2 = sigmoid(np.dot(layer\_1, weights\_1\_2))

vector of

100 cols or

features.

layer\_2\_delta = layer\_2 - y

truth or label  
 $g'(0 \text{ or } 1)$

Backprop  
Sigmoid

Layer\_1\_delta = layer\_2\_delta . dot(weights\_1\_2.T)

Linear weights\_0\_1[x] = layer\_1\_delta \* alpha

softmax weights\_1\_2 = np.outer(layer\_1, layer\_2\_delta) \*  
a output 100 neurons in hidden layer  
alpha

blow 0 to 1

signifying if (np.abs(layer\_2\_delta) < 0.5):

Probability correct += 1

total += 1

if (i % 10 == 9):

progress = str(i / float(len(input\_dataset)))

sys.stdout.write('progress:' + progress[2:4] +

'.' + progress[4:6] +

'%. training Accuracy:' +

+ str(correct / float(total))

+'%')

print()

testing accuracy:

correct, total = (0, 0)

for i in range(len(input\_dataset)-1000, len(input\_data\_set)):

```
x = input_dataset[i] # review  
y = target_dataset[i] # label / actual output
```

```
layer_1 = sigmoid(np.sum(weights_0_1[x], axis=0))  
layer_2 = sigmoid(np.dot(layer_1, weights_1_2))
```

```
# if (np.abs(layer_2 - y) < 0.5):  
if output  
    correct += 1  
prob > 0.5  
    total += 1
```

that means

guessed  
correctly. print ("Test accuracy: " + str(correct/float(total)))

\*) Testing accuracy  $\rightarrow 0.849$  — accurate 84.9%  
accuracy.

\*) Network identifies words, that have either a positive  
negative correlation.

\*) Fundamentally, hidden layers are about grouping  
data points from a previous layer into  $n$  groups  
(where  $n$  — number of neurons in the hidden layer).

Each hidden neuron takes in a datapoint and answers the question "is this datapoint in my group"

\*) An input data point grouping is useful if it does two things. First, the grouping must be useful to the prediction of an output layer.

\*) Second, a grouping is useful if it's an actual phenomenon in the data that you care about.

Good groupings pick up on phenomena that are useful linguistically.

\*) If you can construct two examples with an identical hidden layer, one with the pattern you find interesting and one without, the network is unlikely to find that pattern — thinking of grouping using not or negation. (example).

\*) A hidden layer fundamentally groups the previous layer's data

\* ) At a granular level , each neuron classifies a datapoint as either subscribing or not subscribing to its group .

( movie reviews )

\* ) At a higher level , two data points are similar if they subscribe to many of the same groups .

\* ) Finally , two inputs ( words ) are similar if weights linking them to various hidden neurons ( a measure of each word's group affinity ) are similar .

\* ) Words that have similar predictive power should subscribe to similar groups ( hidden neuron configurations ) .

\* ) Words that correlate with similar labels ( positive or negative ) will have similar weights connecting them to various neurons .

\* ) This phenomenon can be noticed by taking a particular pos or neg word and searching for the

Other word with the most similar weight

Values

\* Words that subscribe to similar groups will have similar predictive power for positive or negative labels. So words that subscribe to similar groups, having similar weight values, will also have similar meaning.

\* To figure out which words are most similar to a target term, you compare each word's vector (row of the matrix) to that of the target term. Using dot product works as well for similarity.

```
from collections import Counter
import math

def similar(target='beautiful'):
    target_index = word2index[target]
    scores = Counter()
    for word, index in word2index.items():
        raw_difference = weights_0_1[index] - (weights_0_1[target_index])
        squared_difference = raw_difference * raw_difference
        scores[word] = -math.sqrt(sum(squared_difference))

    return scores.most_common(10)
```

Euclidean distance.

This allows you to easily query for the most similar word (neuron) according to the network:

print(similar('beautiful'))	print(similar('terrible'))
[('beautiful', -0.0), ('atmosphere', -0.70542101298), ('heart', -0.7339429768542354), ('tight', -0.7470388145765346), ('fascinating', -0.7549291974), ('expecting', -0.759886970744), ('beautifully', -0.7603669338), ('awesome', -0.76647368382398), ('masterpiece', -0.7708280057), ('outstanding', -0.7740642167)]	[('terrible', -0.0), ('dull', -0.760788602671491), ('lacks', -0.76706470275372), ('boring', -0.7682894961694), ('disappointing', -0.768657), ('annoying', -0.78786389931), ('poor', -0.825784172378292), ('horrible', -0.83154121717), ('laughable', -0.8340279599), ('badly', -0.84165373783678)]

Since, the network has only two labels, the input terms are grouped according to which label they tend to predict.

•) input\_size: # of expected features in the input x (depth)

# input of 20 values in sequence and 1 in input size features (depth)

•) hidden\_size: # of features / columns the output of RNN and hidden state will have.

•) batch\_first: if True — input and output tensors will (batch\_size, seq, feature) — Tensor.  
batch\_size as 1st dimension

Seq\_length = 20 # since RNNs work with sequence data.

- Create sample input and target sequence of data points of length 20.

- Sequence Length — # of words in a sentence.

21 values for x-axis all the way up to pi.

time\_steps = np.linspace(0, np.pi, seq\_length+1)

data = np.sin(time\_steps) # y values.

data.resize((seq\_length+1, 1)) # size becomes  
(seq\_length+1, 1), adds an input\_size dimension

x = data[:-1] # input data } 20 datapoints

y = data[1:] # target data } y is basically x  
but shifted

plt.plot(time\_steps[1:], x, 'r.') # input  
plt.plot(time\_steps[1:], y, 'b.') # target  
x target  
y one timestep  
in the future.

Recurrent layer - documentation - responsible for  
calculating a hidden states based on its input.

Class RNN(nn.Module):

def \_\_init\_\_(self, in\_size, out\_size, hid\_dim, n\_lay)

super(RNN, self).\_\_init\_\_()

# vector components  
defines number

self.hidden\_dim = hid\_dim # features the output  
# this RNN  
will have.

self.rnn = nn.RNN(in\_size, hid\_dim, n\_lay, batch\_

first=True,

self.fc = nn.Linear(hid\_dim, out\_size)

↑  
output of # of outputs.

## RNN

hidden state.  
↓

```
def forward(self, x, hidden):
```

```
batch_size = x.size(0)
```

```
r_out, hidden = self.rnn(x, hidden)
```

```
# Shape r_out r_out = r_out.view(-1, self.hidden  
to(batch_size * dimension)
```

$\begin{matrix} \text{seq\_length, hidden\_dim} \\ \uparrow \\ \text{cols} \end{matrix}$

```
output = self.fc(r_out)
```

```
return output, hidden
```

batch seq-1  
input - [ 1, 20, 1 ]  
feature

output - [ 20, 1 ]  
Batch output  
size

$(20, 10)$

Hidden -  $[2, 10, 1]$   
 layers      dimension

$$(1, 1) \circ (1 \circ 32) \rightarrow (1, 32)$$

hidden state  
output neurons

$$r_{out} = (20, 32) \Rightarrow 20 r's$$

32 C's

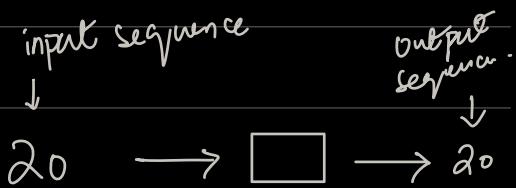
$$w \rightarrow (32, 1)$$

$$I - (20, 32)$$

$$w \rightarrow (32, 1)$$

where every column is a value for a predicted sequence.

\* For every iteration or a sequence, a non-hint added.



- Deep Learning framework provides support for automatic backpropagation and automatic optimization. These features let you specify only the forward propagation code of a model with the framework taking care of backpropagation and weight updates automatically.
- Providing high level interfaces to common layers and loss functions.